

▼ Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split, cross_val_score
import plotly.express as px
import plotly.graph_objects as go
```

```
houses_data = pd.read_csv("train.csv")
test_data = pd.read_csv("test.csv")
```

▼ Data Cleaning and Preparation

```
# Display basic information about the dataset
print("Step 2: Data Collection")
print("\nDataset loaded successfully.")
print("\nBasic Information About the Dataset:")
print("Number of Rows:", len(houses_data))
print("Number of Columns:", len(houses_data.columns))
print("\nSample Data (first 5 rows):")
print(houses_data.head())
```

↗ Step 2: Data Collection

Dataset loaded successfully.

Basic Information About the Dataset:

Number of Rows: 1460

Number of Columns: 81

Sample Data (first 5 rows):

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	\
0	Lv1	AllPub	...	0	NaN	NaN	NaN	0	2	
1	Lv1	AllPub	...	0	NaN	NaN	NaN	0	5	
2	Lv1	AllPub	...	0	NaN	NaN	NaN	0	9	
3	Lv1	AllPub	...	0	NaN	NaN	NaN	0	2	
4	Lv1	AllPub	...	0	NaN	NaN	NaN	0	12	

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

```
houses_data.info()
```



```
38 TotalBsmntSF      1460 non-null    int64
39 Heating           1460 non-null    object
40 HeatingQC          1460 non-null    object
41 CentralAir         1460 non-null    object
42 Electrical         1459 non-null    object
43 1stFlrSF           1460 non-null    int64
44 2ndFlrSF           1460 non-null    int64
45 LowQualFinSF       1460 non-null    int64
46 GrLivArea          1460 non-null    int64
47 BsmtFullBath       1460 non-null    int64
48 BsmtHalfBath       1460 non-null    int64
49 FullBath           1460 non-null    int64
50 HalfBath           1460 non-null    int64
51 BedroomAbvGr      1460 non-null    int64
52 KitchenAbvGr       1460 non-null    int64
53 KitchenQual        1460 non-null    object
54 TotRmsAbvGrd      1460 non-null    int64
55 Functional         1460 non-null    object
56 Fireplaces         1460 non-null    int64
57 FireplaceQu       770 non-null     object
58 GarageType         1379 non-null    object
59 GarageYrBlt        1379 non-null    float64
60 GarageFinish       1379 non-null    object
61 GarageCars         1460 non-null    int64
62 GarageArea         1460 non-null    int64
63 GarageQual         1379 non-null    object
64 GarageCond         1379 non-null    object
65 PavedDrive         1460 non-null    object
66 WoodDeckSF         1460 non-null    int64
67 OpenPorchSF        1460 non-null    int64
68 EnclosedPorch      1460 non-null    int64
69 3SsnPorch          1460 non-null    int64
70 ScreenPorch        1460 non-null    int64
71 PoolArea           1460 non-null    int64
72 PoolQC             7 non-null      object
73 Fence              281 non-null    object
74 MiscFeature        54 non-null     object
75 MiscVal            1460 non-null    int64
76 MoSold             1460 non-null    int64
77 YrSold             1460 non-null    int64
78 SaleType           1460 non-null    object
79 SaleCondition       1460 non-null    object
80 SalePrice          1460 non-null    int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB
```


```
houses_data.describe()
```



	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	\
count	1460.000000	1460.000000	1201.000000	1460.000000	1460.000000	1460.000000	14
mean	730.500000	56.897260	70.049958	10516.828082	6.099315	5.575342	19
std	421.610009	42.300571	24.284752	9981.264932	1.382997	1.112799	
min	1.000000	20.000000	21.000000	1300.000000	1.000000	1.000000	18
25%	365.750000	20.000000	59.000000	7553.500000	5.000000	5.000000	19
50%	730.500000	50.000000	69.000000	9478.500000	6.000000	5.000000	19
75%	1095.250000	70.000000	80.000000	11601.500000	7.000000	6.000000	20
max	1460.000000	190.000000	313.000000	215245.000000	10.000000	9.000000	20

8 rows × 8 columns

```
houses_data.isnull().sum()
```



```
Id      0
MSSubClass  0
MSZoning  0
LotFrontage  259
LotArea    0
...
MoSold    0
YrSold    0
SaleType  0
SaleCondition  0
SalePrice  0
Length: 81, dtype: int64
```

```
# Extract columns with null values
columns_with_null = houses_data.columns[houses_data.isnull().any()]
```

```
# Display columns with null values
print("Columns with null values:")
for col in columns_with_null:
    print(col)
```

```
Columns with null values:
LotFrontage
Alley
MasVnrType
MasVnrArea
BsmtQual
BsmtCond
BsmtExposure
BsmtFinType1
BsmtFinType2
Electrical
FireplaceQu
GarageType
GarageYrBlt
GarageFinish
GarageQual
GarageCond
PoolQC
Fence
MiscFeature
```

```
# Check for missing values in each column
missing_values = houses_data.isnull().sum()
```

```
# Print columns with missing values and their corresponding counts
columns_with_missing_values = missing_values[missing_values > 0]
print("\nColumns with Missing Values:")
print(columns_with_missing_values)
```

```
Columns with Missing Values:
LotFrontage      259
Alley            1369
MasVnrType        872
MasVnrArea         8
BsmtQual          37
BsmtCond          37
BsmtExposure      38
BsmtFinType1      37
BsmtFinType2      38
Electrical         1
FireplaceQu       690
GarageType         81
GarageYrBlt        81
GarageFinish       81
GarageQual         81
GarageCond         81
PoolQC            1453
Fence             1179
MiscFeature       1406
dtype: int64
```

```
# Check for duplicate rows
duplicates_before = houses_data.duplicated().sum()
```

```
# Remove duplicate rows
houses_data.drop_duplicates(inplace=True)
```

```
# Check for duplicate rows after removal
duplicates_after = houses_data.duplicated().sum()
```

```
# Print the results
if duplicates_before > 0:
    print(f"Handling Duplicates\n{duplicates_before} duplicate row(s) were found and removed.")
else:
    print("Handling Duplicates\nNo duplicate rows found in the dataset.")
```

```
Handling Duplicates
No duplicate rows found in the dataset.
```

```
# Get the column names and data types
column_info = houses_data.dtypes


# Display column names and data types horizontally
for col_name, data_type in column_info.items(): # Use items() instead of iteritems()
    print(f'{col_name}: {data_type}\t", end='')
```

```
Id: int64      MSSubClass: int64      MSZoning: object      LotFrontage: float64      LotArea: int64      Street: object      Alley: object
```

```
houses_data.columns
```

```
Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
       'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
       'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
       'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
       'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
       'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
       'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
       'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
       'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
       'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
       'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
       'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
       'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
       'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
       'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
       'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
       'SaleCondition', 'SalePrice'],
      dtype='object')
```

```
fig = px.scatter_3d(houses_data, x = 'LotArea', y='LotFrontage', z = 'SalePrice')
fig.show()
```




```
# list of categorical columns to exclude
categorical_columns = ['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope', 'Neighborhood', 'Con

# Select numerical columns
numerical_columns = [col for col in houses_data.columns if col not in categorical_columns]

# Create a DataFrame with only numerical features and the target variable
numerical_data = houses_data[numerical_columns + ['SalePrice']]

# Calculate the correlation matrix
correlation_matrix = numerical_data.corr()

# Step 2: Generate a heatmap
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix[['SalePrice']], annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Feature Correlation with SalePrice")
plt.show()
```




```
# Split the data into training and testing sets
X = houses_data[['TotalBsmtSF', 'BedroomAbvGr', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath']]
y = houses_data['SalePrice']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Select the columns of interest (features and target)
features = houses_data[['TotalBsmtSF', 'BedroomAbvGr', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath']]
target = houses_data[['SalePrice']]

# Create a new DataFrame with only the selected columns
data_subset = pd.concat([features, target], axis=1) # Use square brackets and specify axis=1

# Calculate the correlation matrix
correlation_matrix = data_subset.corr()

# Create a heatmap for the correlation matrix
plt.figure(figsize=(8, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='Blues', fmt=".2f")
plt.title("Correlation Heatmap: Features vs. Target")
plt.show()
```



```
# Check for missing values in the training dataset
missing_train = houses_data.isnull().sum()
print("Missing Values in Training Data:")
print(missing_train)

# Check for missing values in the testing dataset
missing_test = test_data.isnull().sum()
print("\nMissing Values in Testing Data:")
print(missing_test)
```



```
Missing Values in Training Data:
Id                0
MSSubClass        0
MSZoning          0
LotFrontage      259
LotArea           0
...
MoSold           0
```

```
YrSold      0
SaleType    0
SaleCondition 0
SalePrice   0
Length: 81, dtype: int64

Missing Values in Testing Data:
Id          0
MSSubClass  0
MSZoning    4
LotFrontage 227
LotArea     0
...
MiscVal     0
MoSold      0
YrSold      0
SaleType    1
SaleCondition 0
Length: 80, dtype: int64
```

```
# Missing values in selected features
features_missing_values = X.isnull().sum()
features_missing_values
```

```
TotalBsmtSF      0
BedroomAbvGr     0
BsmtFullBath     0
BsmtHalfBath     0
FullBath         0
HalfBath         0
dtype: int64
```

```
#create a linear regression model
model = LinearRegression()
```

```
# Fit the model to the training data
model.fit(X_train, y_train)
print(model)
```

```
LinearRegression()
```

```
# Make predictions on the test data
y_pred = model.predict(X_test)
```

```
# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

```
print(f"Mean Squared Error: {mse:.2f}")
print(f"R-squared: {r2:.2f}")
```

```
Mean Squared Error: 2693215363.50
R-squared: 0.65
```

```
X.sample(5)
```

	TotalBsmtSF	BedroomAbvGr	BsmtFullBath	BsmtHalfBath	FullBath	HalfBath
1024	1565	2	1	0	2	0
607	896	3	1	0	3	0
1176	876	3	1	0	1	0
1428	788	2	1	0	1	0
914	612	2	0	0	2	1

```
# Predict the price of a new house
new_house = np.array([[2500, 3, 1,0,2,1]])
predicted_price = model.predict(new_house)
print(f"Predicted Price for the New House: ${predicted_price[0]:.2f}")
```

```
Predicted Price for the New House: $354103.01
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning:
```

X does not have valid feature names, but LinearRegression was fitted with feature names

```
# Cross-validation to assess model performance
cv_scores = cross_val_score(model, X, y, cv=5)
print('Cross-Validation Scores:', cv_scores)
print('Mean CV Score:', cv_scores.mean())
```

```
↗ Cross-Validation Scores: [0.60009438 0.65556307 0.63763301 0.61946238 0.44203474]
Mean CV Score: 0.5909575152454488
```

```
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Prices vs. Predicted Prices")
plt.show()
```

